

Service Template Library

Service Template Library [🔗](#)

Overview [🔗](#)

This document catalogs all available service templates that can be used with the integrated cluster-snek platform. Each service includes configuration options, dependencies, and example usage.

Core Services [🔗](#)

Authentication & Identity [🔗](#)

Keycloak SSO [🔗](#)

Single Sign-On and identity management platform.

```
1 services:
2   keycloak:
3     enabled: true
4     version: "23.0.0"
5     replicas: 2
6     database:
7       type: "postgres"
8       size: "10Gi"
9     features:
10      - user-federation
11      - social-login
12      - mfa
13     integrations:
14      - gitlab
15      - argocd
16      - grafana
17      - jupyterhub
```

Dependencies: PostgreSQL database

Resource Requirements: 2 CPU, 4GB RAM minimum

Source Control [🔗](#)

GitLab CE [🔗](#)

Complete DevOps platform with Git repository management.

```
1 services:
2   gitlab:
3     enabled: true
4     version: "16.6.0"
5     storage:
6       repos: "100Gi"
7       lfs: "50Gi"
8       artifacts: "50Gi"
9       uploads: "10Gi"
10    features:
11      - ci-cd
```

```

12     - container-registry
13     - package-registry
14     - pages
15     runners:
16       count: 3
17       concurrent: 10

```

Dependencies: PostgreSQL, Redis, Object Storage (optional)

Resource Requirements: 4 CPU, 8GB RAM minimum

AI/ML Platform [🔗](#)

JupyterHub [🔗](#)

Multi-user notebook environment for data science and AI/ML workloads.

```

1  services:
2    jupyterhub:
3      enabled: true
4      version: "3.1.0"
5      authentication:
6        type: "keycloak"
7      profiles:
8        - name: "minimal"
9          cpu: 1
10         memory: "2Gi"
11        - name: "standard"
12          cpu: 2
13          memory: "4Gi"
14        - name: "gpu"
15          cpu: 4
16          memory: "8Gi"
17          gpu: 1
18      storage:
19        home_size: "10Gi"
20        shared_size: "100Gi"

```

Dependencies: Keycloak (for SSO)

Resource Requirements: Variable based on user profiles

Vector Databases [🔗](#)

Weaviate [🔗](#)

Production-grade vector database for AI applications.

```

1  services:
2    weaviate:
3      enabled: true
4      version: "1.24.0"
5      replicas: 3
6      modules:
7        - text2vec-transformers
8        - generative-openai
9        - qna-transformers
10     storage:
11       size: "100Gi"
12       class: "fast-ssd"
13     resources:

```

```
14   cpu: 4
15   memory: "16Gi"
```

Dependencies: None

Resource Requirements: 4 CPU, 16GB RAM per replica

Monitoring & Observability [🔗](#)

Prometheus Stack [🔗](#)

Complete monitoring solution with Prometheus, Grafana, and AlertManager.

```
1  services:
2    monitoring:
3      enabled: true
4      components:
5        prometheus:
6          retention: "30d"
7          storage: "100Gi"
8        grafana:
9          dashboards:
10             - kubernetes
11             - argocd
12             - ai-workloads
13             - security
14        alertmanager:
15          receivers:
16            - slack
17            - email
18        loki:
19          retention: "14d"
20          storage: "50Gi"
```

Dependencies: None

Resource Requirements: 2 CPU, 4GB RAM for full stack

Backup & Recovery [🔗](#)

Velero [🔗](#)

Kubernetes backup and disaster recovery solution.

```
1  services:
2    velero:
3      enabled: true
4      version: "1.12.0"
5      provider: "aws" # or azure, gcp, minio
6      backup:
7        schedule: "0 2 * * *" # Daily at 2 AM
8        retention: "30d"
9        include_namespaces:
10           - default
11           - applications
12           - ai-workloads
13        exclude_resources:
14           - events
15           - pods/log
16      storage:
17        bucket: "velero-backups"
```

```
18     region: "us-east-1"
```

Dependencies: Object storage (S3 compatible)

Resource Requirements: 1 CPU, 2GB RAM

Security & Compliance [🔗](#)

Cerbos [🔗](#)

Authorization engine for fine-grained access control.

```
1  services:
2    cerbos:
3      enabled: true
4      version: "0.30.0"
5      replicas: 3
6      policy_repository:
7        url: "https://github.com/org/cerbos-policies"
8        branch: "main"
9      audit:
10       enabled: true
11       backend: "postgres"
12       retention: "1y"
13     integrations:
14       - weaviate
15       - gitlab
16       - jupyterhub
```

Dependencies: PostgreSQL (for audit logs)

Resource Requirements: 1 CPU, 2GB RAM per replica

Service Profiles [🔗](#)

Development Profile [🔗](#)

Minimal resources for development environments.

```
1  service_profile: "development"
2  services:
3    gitlab:
4      resources_override:
5        cpu: 2
6        memory: "4Gi"
7    monitoring:
8      components:
9        - prometheus
10       - grafana
11    minimal: true
```

Production Profile [🔗](#)

High availability and full features.

```
1  service_profile: "production"
2  services:
3    keycloak:
4      replicas: 3
5      ha: true
6    gitlab:
```

```

7     replicas: 2
8     ha: true
9     monitoring:
10    ha: true
11    long_term_storage: true

```

Service Dependencies [🔗](#)

Dependency Graph [🔗](#)

```

1 keycloak
2   └─ postgresql
3 gitlab
4   └─ postgresql
5   └─ redis
6   └─ keycloak (for SS0)
7 jupyterhub
8   └─ keycloak (for SS0)
9 monitoring
10  └─ (independent)
11 cerbos
12  └─ postgresql (for audit)
13 velero
14  └─ object storage

```

Automatic Dependency Resolution [🔗](#)

The platform automatically installs dependencies:

```

1 services:
2   - gitlab # Automatically includes postgresql, redis
3   - jupyterhub # Automatically includes keycloak

```

Configuration Examples [🔗](#)

Minimal AI Research Lab [🔗](#)

```

1 clusters:
2   - name: "research"
3     services:
4       - jupyterhub
5       - weaviate
6       - monitoring

```

Full DevOps Platform [🔗](#)

```

1 clusters:
2   - name: "devops"
3     services:
4       - keycloak
5       - gitlab
6       - monitoring
7       - velero
8     service_overrides:
9       gitlab:
10        runners:

```

11 `count: 5`

Secure Enterprise Platform [🔗](#)

```
1 clusters:
2   - name: "enterprise"
3     services:
4       - keycloak
5       - gitlab
6       - jupyterhub
7       - weaviate
8       - cerbos
9       - monitoring
10      - velero
11    security_profile: "zero-trust"
```

Service Customization [🔗](#)

Override Values [🔗](#)

```
1 service_overrides:
2   gitlab:
3     version: "16.7.0" # Override default version
4     custom_values:
5       gitlab:
6         webservice:
7           minReplicas: 3
```

Custom Services [🔗](#)

Add your own service definitions:

```
1 custom_services:
2   my-app:
3     chart:
4       repository: "https://charts.example.com"
5       name: "my-app"
6       version: "1.0.0"
7     values:
8       replicas: 2
9       ingress:
10        enabled: true
```

Best Practices [🔗](#)

1. **Start Small:** Begin with core services and add more as needed
2. **Monitor Resources:** Use monitoring to track resource usage
3. **Security First:** Enable Cerbos for production workloads
4. **Regular Backups:** Always enable Velero for critical data
5. **SSO Everything:** Use Keycloak for centralized authentication